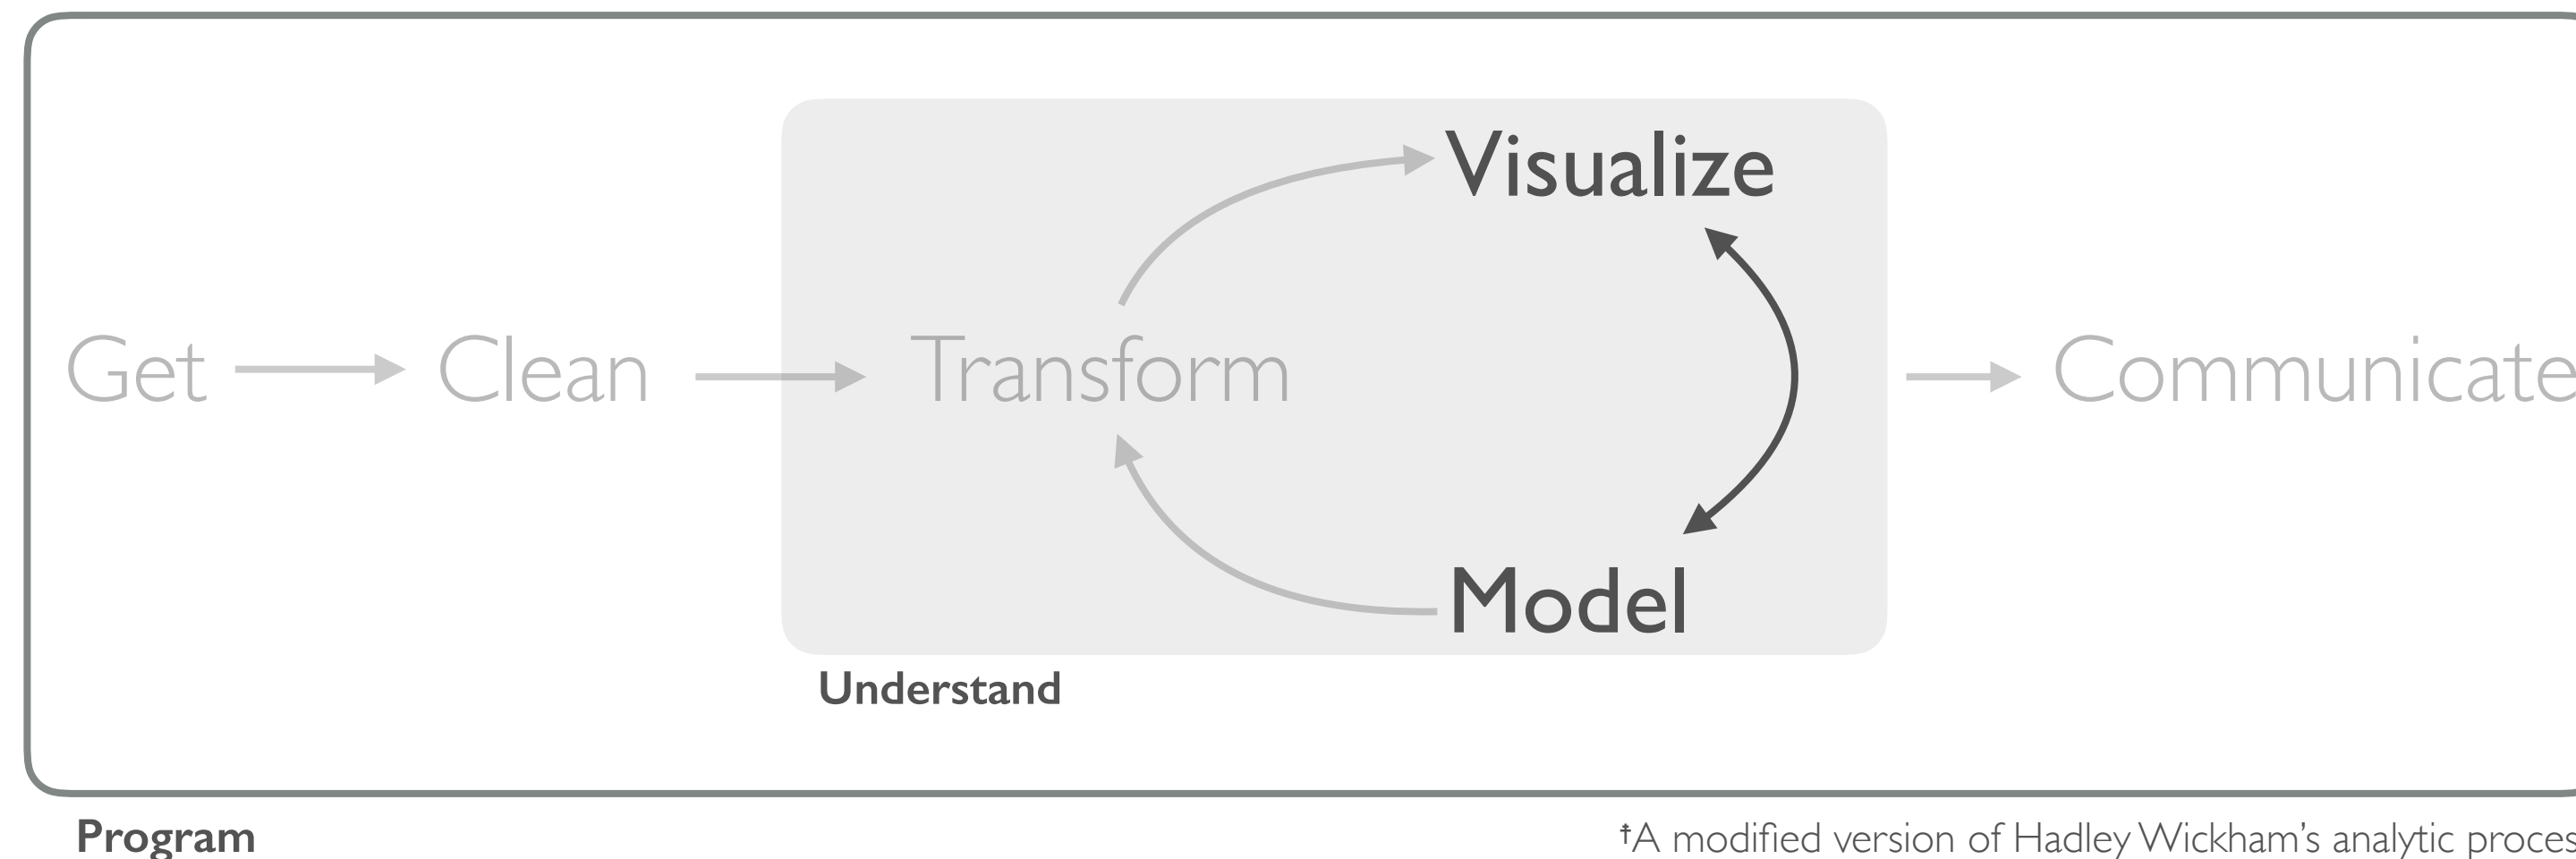



# MANAGING MANY MODELS



# PREREQUISITES



# PREREQUISITES

```
library(modelr)
library(tidyverse)
library(gapminder)
```

```
gapminder
```

```
# A tibble: 1,704 × 6
```

	country	continent	year	lifeExp	pop	gdpPercap
	<fctr>	<fctr>	<int>	<dbl>	<int>	<dbl>
1	Afghanistan	Asia	1952	28.801	8425333	779.4453
2	Afghanistan	Asia	1957	30.332	9240934	820.8530
3	Afghanistan	Asia	1962	31.997	10267083	853.1007
4	Afghanistan	Asia	1967	34.020	11537966	836.1971
5	Afghanistan	Asia	1972	36.088	13079460	739.9811
6	Afghanistan	Asia	1977	38.438	14880372	786.1134
7	Afghanistan	Asia	1982	39.854	12881816	978.0114


HOW DOES LIFE EXPECTANCY CHANGE  
OVER TIME?



# A COMMON TREND...FOR THE MOST PART

```
gapminder %>%  
  ggplot(aes(year, lifeExp, group = country)) +  
  geom_line(alpha = 1/3)
```

We see a fairly common trend  
across most countries



# MODEL THE WHOLE THING

```
full_mod <- lm(lifeExp ~ year + country,  
              data = gapminder)  
  
summary(full_mod)
```

One approach to model this relationship over time is to use the following multivariate regression model


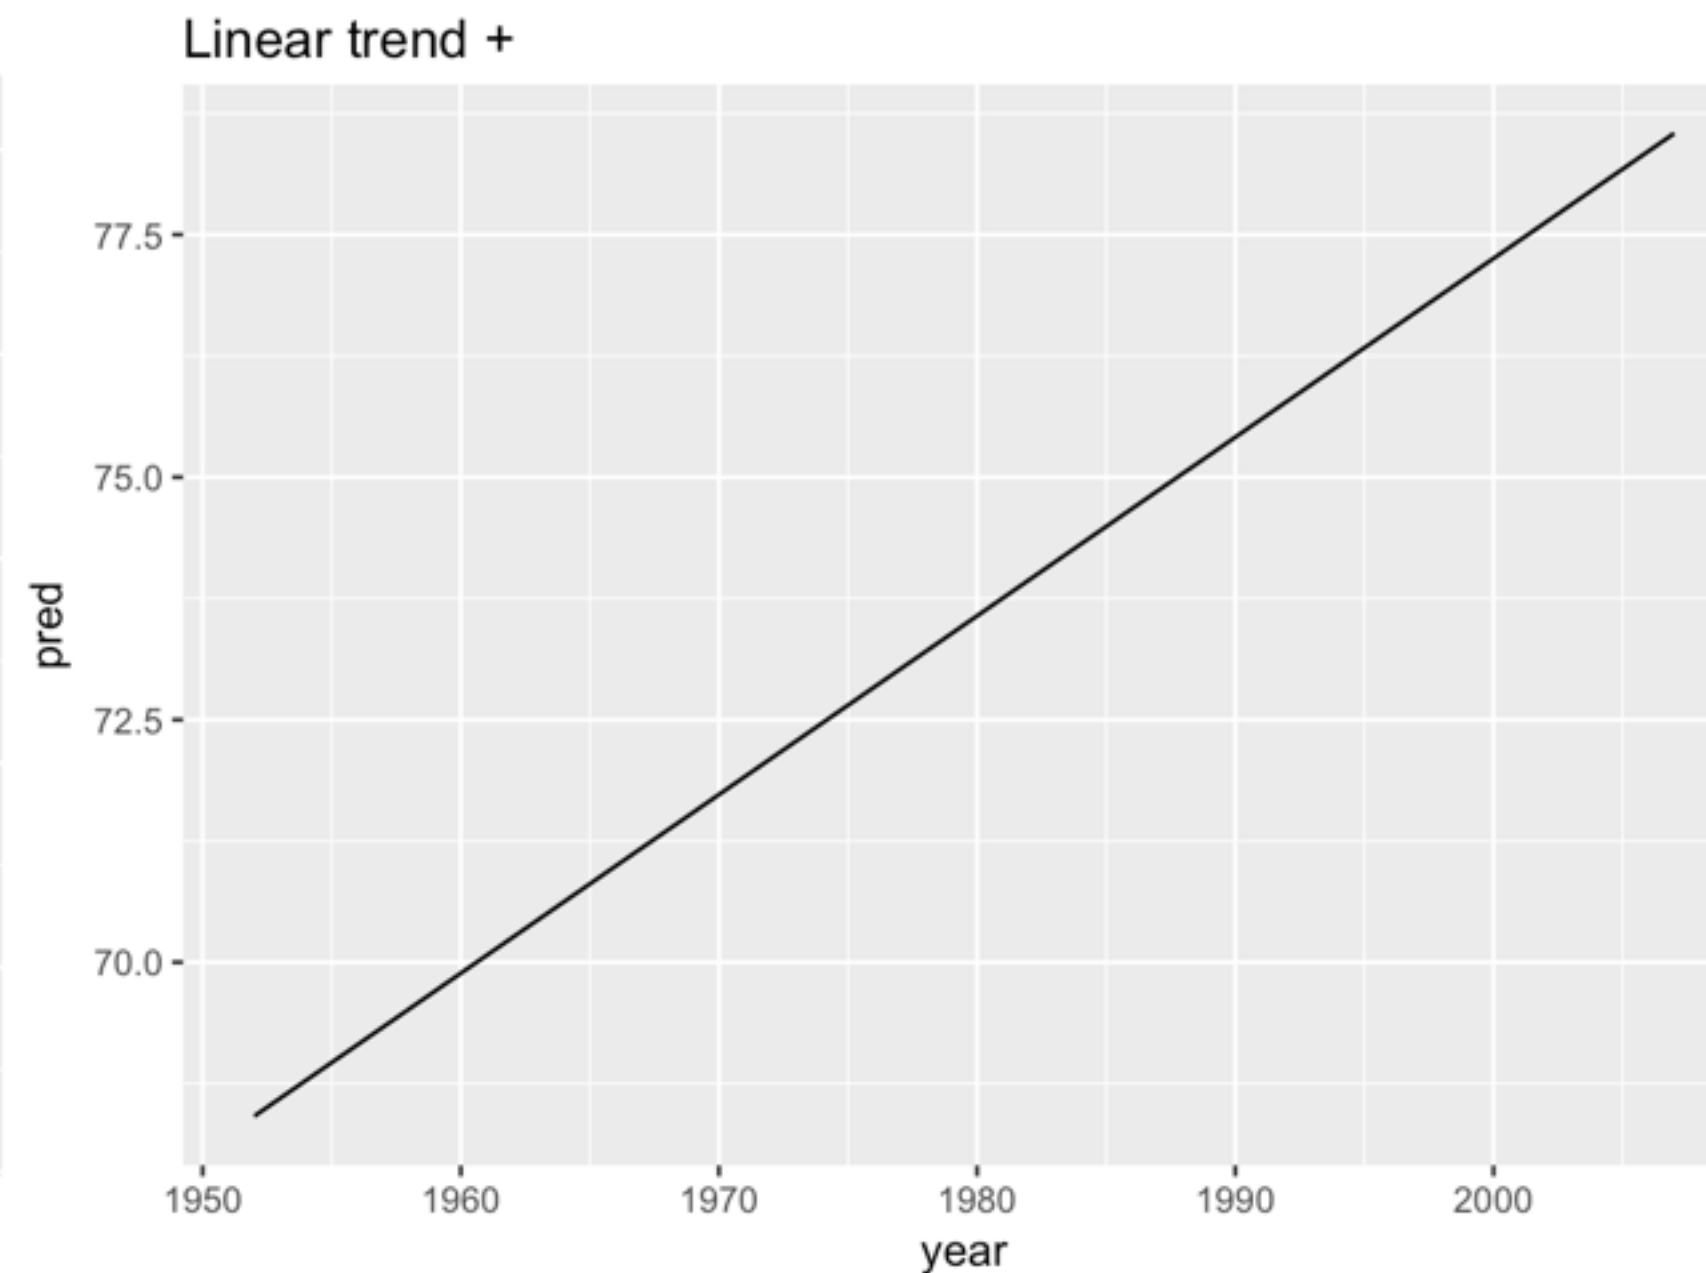
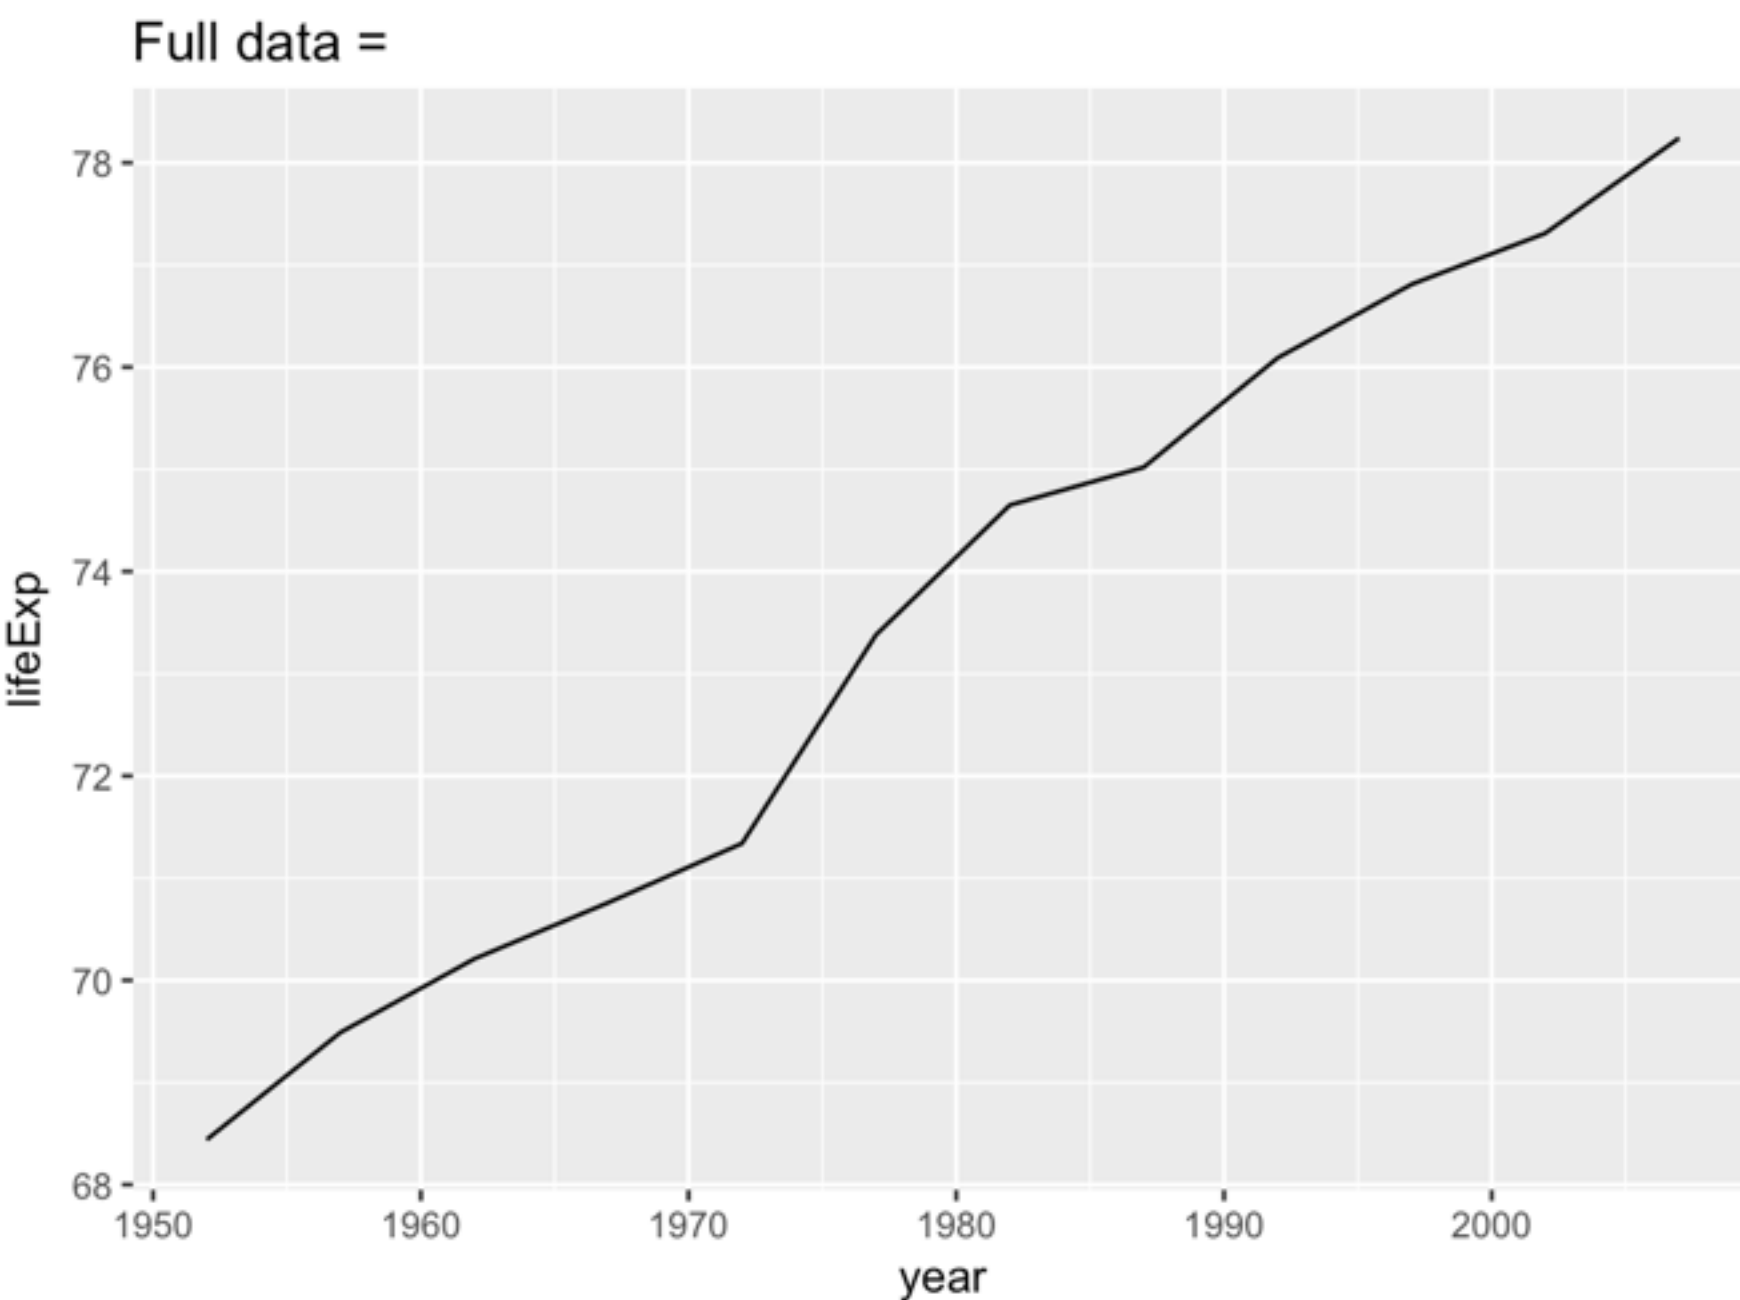
Numerically, it appears to fit very well but...

Visually, there are residual concerns in the latter year predictions

# MODEL A SINGLE PART

```
usa <- filter(gapminder, country == "United States")  
usa_mod <- lm(lifeExp ~ year, data = usa)
```

Alternatively, we may want to just focus on a single country to see the particular relationship there



*But what if we want to compare this country-level model across  
**all** the countries?*



NESTED DATA



# CONFOUNDING VARIABLE

```
by_country <- gapminder %>%  
  group_by(country, continent) %>%  
  nest()
```

```
by_country
```

```
# A tibble: 142 × 3
```

	country	continent	data
	<fctr>	<fctr>	<list>
1	Afghanistan	Asia	<tibble [12 × 4]>
2	Albania	Europe	<tibble [12 × 4]>
3	Algeria	Africa	<tibble [12 × 4]>
4	Angola	Africa	<tibble [12 × 4]>
5	Argentina	Americas	<tibble [12 × 4]>
6	Australia	Oceania	<tibble [12 × 4]>
7	Austria	Europe	<tibble [12 × 4]>

Introducing a new data structure - the **nested** data frame

*What is in each data column element? Can you figure out how to look at this data?*

# CONFOUNDING VARIABLE

```
by_country <- gapminder %>%  
  group_by(country, continent) %>%  
  nest()
```

```
by_country$data[[1]]
```

```
# A tibble: 12 × 4
```

	year	lifeExp	pop	gdpPercap
	<int>	<dbl>	<int>	<dbl>
1	1952	28.801	8425333	779.4453
2	1957	30.332	9240934	820.8530
3	1962	31.997	10267083	853.1007
4	1967	34.020	11537966	836.1971
5	1972	36.088	13079460	739.9811
6	1977	38.438	14880372	786.1134
7	1982	39.854	12881816	978.0114

Called: “list-columns”

- each element is a list
- interact with these elements just like you do a list

# YOUR TURN!

*Discuss with your neighbor how you could use this data structure (along with previously reviewed iterative functions) to apply a country-level model across each of the list-columns.*



# ITERATIVE MODEL APPLICATION

# LET'S DEVELOP A MODEL FUNCTION

```
country_model <- function(df) {  
  lm(lifeExp ~ year, data = df)  
}
```

# LET'S DEVELOP A MODEL FUNCTION

```
country_model <- function(df) {  
  lm(lifeExp ~ year, data = df)  
}  
map(by_country$data, country_model)  
[[1]]
```

Call:

```
lm(formula = lifeExp ~ year, data = df)
```

Coefficients:

(Intercept)	year
-507.5343	0.2753

Remember the **map** function?

We can apply this model over every element in the **data** column with **map**

*How could we add this information to our nested data frame?*

# LET'S DEVELOP A MODEL FUNCTION

```
country_model <- function(df) {  
  lm(lifeExp ~ year, data = df)  
}  
  
by_country <- by_country %>%  
  mutate(model = map(data, country_model))  
  
by_country  
# A tibble: 142 × 4  
   country continent      data      model  
   <fctr>    <fctr>    <list>   <list>  
1 Afghanistan    Asia <tibble [12 × 4]> <S3: lm>  
2  Albania    Europe <tibble [12 × 4]> <S3: lm>  
3  Algeria    Africa <tibble [12 × 4]> <S3: lm>  
4   Angola    Africa <tibble [12 × 4]> <S3: lm>
```

Using **mutate** will save these regression results in a new **list-column model** variable

Now we have all our model results neatly packaged together with each country



UNNESTING



# GETTING STUFF OUT OF OUR NEST

```
by_country %>%
  mutate(resids = map2(data, model, add_residuals))
```

	country	continent	data	model	resids
	<fctr>	<fctr>	<list>	<list>	<list>
1	Afghanistan	Asia	<tibble [12 × 4]>	<S3: lm>	<tibble [12 × 5]>
2	Albania	Europe	<tibble [12 × 4]>	<S3: lm>	<tibble [12 × 5]>
3	Algeria	Africa	<tibble [12 × 4]>	<S3: lm>	<tibble [12 × 5]>
4	Angola	Africa	<tibble [12 × 4]>	<S3: lm>	<tibble [12 × 5]>
5	Argentina	Americas	<tibble [12 × 4]>	<S3: lm>	<tibble [12 × 5]>
6	Australia	Oceania	<tibble [12 × 4]>	<S3: lm>	<tibble [12 × 5]>
7	Austria	Europe	<tibble [12 × 4]>	<S3: lm>	<tibble [12 × 5]>
8	Bahrain	Asia	<tibble [12 × 4]>	<S3: lm>	<tibble [12 × 5]>
9	Bangladesh	Asia	<tibble [12 × 4]>	<S3: lm>	<tibble [12 × 5]>
10	Belgium	Europe	<tibble [12 × 4]>	<S3: lm>	<tibble [12 × 5]>
#	... with 132 more rows				

## Using map2

- Similar to the **map** function
- But uses two argument inputs to map a function over (i.e. data, model)

*Can you figure out what this code is doing?*

# GETTING STUFF OUT OF OUR NEST

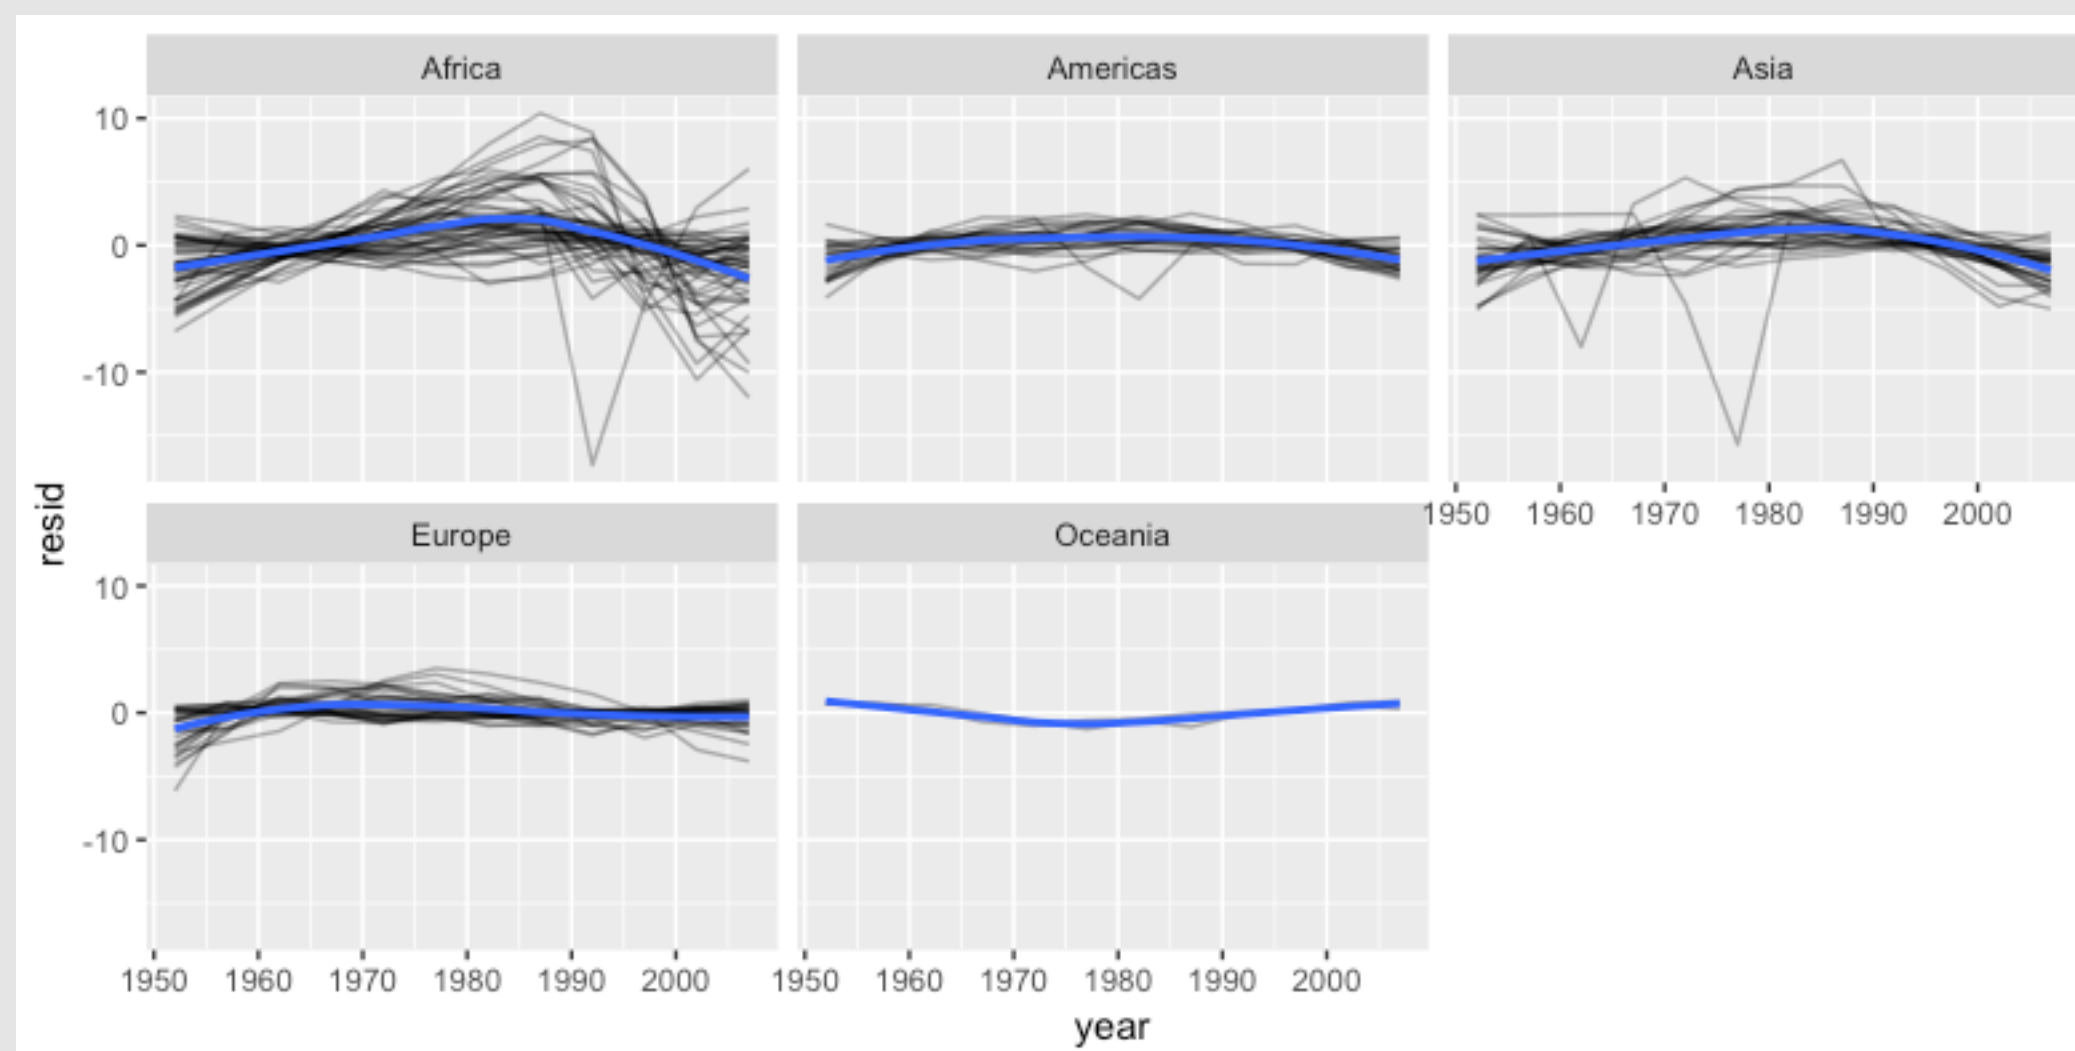
```
by_country %>%
  mutate(resids = map2(data, model, add_residuals)) %>%
  unnest(resids)
# A tibble: 1,704 × 7
   country continent  year lifeExp      pop gdpPercap      resid
   <fctr>    <fctr> <int>   <dbl>    <int>    <dbl>    <dbl>
1 Afghanistan Asia  1952  28.801  8425333  779.4453 -1.10629487
2 Afghanistan Asia  1957  30.332  9240934  820.8530 -0.95193823
3 Afghanistan Asia  1962  31.997 10267083  853.1007 -0.66358159
4 Afghanistan Asia  1967  34.020 11537966  836.1971 -0.01722494
5 Afghanistan Asia  1972  36.088 13079460  739.9811  0.67413170
6 Afghanistan Asia  1977  38.438 14880372  786.1134  1.64748834
7 Afghanistan Asia  1982  39.854 12881816  978.0114  1.68684499
8 Afghanistan Asia  1987  40.822 13867957  852.3959  1.27820163
9 Afghanistan Asia  1992  41.674 16317921  649.3414  0.75355828
10 Afghanistan Asia  1997  41.763 22227415  635.3414 -0.53408508
# ... with 1,694 more rows
```

...but our data is still nested

We can use **unnest** to extract the nested information of choice and convert back to a regular data frame

# GETTING STUFF OUT OF OUR NEST

```
by_country %>%  
  mutate(resids = map2(data, model, add_residuals)) %>%  
  unnest(resids) %>%  
  ggplot(aes(year, resid)) +  
    geom_line(aes(group = country), alpha = 1 / 3) +  
    geom_smooth(se = FALSE) +  
    facet_wrap(~ continent)
```



This allows us to flow right into our normal visualization of residuals to compare across continents and countries

MODEL QUALITY



# ASSESSING OUR MODEL(S)

```
usa <- filter(gapminder, country == "United States")  
usa_mod <- lm(lifeExp ~ year, data = usa)
```

Remember our single country model?

*How would you normally assess  
model performance  
(numerically)?*

# ASSESSING OUR MODEL(S)

```
usa <- filter(gapminder, country == "United States")
usa_mod <- lm(lifeExp ~ year, data = usa)
summary(usa_mod)
Call:
lm(formula = lifeExp ~ year, data = usa)
```

Residuals:

Min	1Q	Median	3Q	Max
-0.75723	-0.30394	0.06735	0.19752	0.71108

Coefficients:

	Estimate	Std. Error	t value	Pr(> t )
(Intercept)	-291.08449	13.77740	-21.13	1.25e-09 ***
year	0.18417	0.00696	26.46	1.37e-10 ***

---

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

Remember our single country model?

**summary()**



# ASSESSING OUR MODEL(S)

```
usa <- filter(gapminder, country == "United States")
usa_mod <- lm(lifeExp ~ year, data = usa)
broom::glance(usa_mod)
# A tibble: 1 × 11
#   r.squared adj.r.squared sigma statistic
#   <dbl>      <dbl>      <dbl>      <dbl>
1 0.9859202    0.9845122 0.4161339  700.2351
# ... with 7 more variables: p.value <dbl>, df <int>,
#   logLik <dbl>, AIC <dbl>, BIC <dbl>,
#   deviance <dbl>, df.residual <int>
```

An alternative approach is with  
**broom::glance**

Creates a tidy one-row data  
frame with useful model results



# ASSESSING OUR MODEL(S)

```
by_country
# A tibble: 142 × 4
  country continent      data      model
  <fctr>    <fctr>      <list>    <list>
1 Afghanistan      Asia <tibble [12 × 4]> <S3: lm>
2      Albania    Europe <tibble [12 × 4]> <S3: lm>
3      Algeria    Africa <tibble [12 × 4]> <S3: lm>
4      Angola    Africa <tibble [12 × 4]> <S3: lm>
5    Argentina Americas <tibble [12 × 4]> <S3: lm>
6    Australia Oceania <tibble [12 × 4]> <S3: lm>
7      Austria    Europe <tibble [12 × 4]> <S3: lm>
8      Bahrain      Asia <tibble [12 × 4]> <S3: lm>
9 Bangladesh      Asia <tibble [12 × 4]> <S3: lm>
10     Belgium    Europe <tibble [12 × 4]> <S3: lm>
# with 132 more rows
```

*How could we use this with our many models approach?*

# ASSESSING OUR MODEL(S)

```
by_country %>%
```

```
  mutate(glance = map(model, broom::glance))
```

```
# A tibble: 142 × 5
```

	country	continent		data	model	glance
	<fctr>	<fctr>		<list>	<list>	<list>
1	Afghanistan	Asia	<tibble [12 × 4]>	<S3: lm>	<data.frame [1 × 11]>	
2	Albania	Europe	<tibble [12 × 4]>	<S3: lm>	<data.frame [1 × 11]>	
3	Algeria	Africa	<tibble [12 × 4]>	<S3: lm>	<data.frame [1 × 11]>	
4	Angola	Africa	<tibble [12 × 4]>	<S3: lm>	<data.frame [1 × 11]>	
5	Argentina	Americas	<tibble [12 × 4]>	<S3: lm>	<data.frame [1 × 11]>	
6	Australia	Oceania	<tibble [12 × 4]>	<S3: lm>	<data.frame [1 × 11]>	
7	Austria	Europe	<tibble [12 × 4]>	<S3: lm>	<data.frame [1 × 11]>	
8	Bahrain	Asia	<tibble [12 × 4]>	<S3: lm>	<data.frame [1 × 11]>	
9	Bangladesh	Asia	<tibble [12 × 4]>	<S3: lm>	<data.frame [1 × 11]>	
10	Belgium	Europe	<tibble [12 × 4]>	<S3: lm>	<data.frame [1 × 11]>	

# ASSESSING OUR MODEL(S)

```
by_country %>%  
  mutate(glance = map(model, broom::glance)) %>%  
  unnest(glance, .drop = TRUE)  
# A tibble: 142 × 13  
  country continent r.squared adj.r.squared sigma statistic p.value  
  <fctr>    <fctr>    <dbl>         <dbl>    <dbl>    <dbl>    <dbl>  
1  Afghanistan      Asia 0.9477123      0.9424835 1.2227880 181.24941 9.835213e-08  
2    Albania      Europe 0.9105778      0.9016355 1.9830615 101.82901 1.462763e-06  
3    Algeria      Africa 0.9851172      0.9836289 1.3230064 661.91709 1.808143e-10  
4     Angola      Africa 0.8878146      0.8765961 1.4070091 79.13818 4.593498e-06
```

*This allows us to quickly assess and compare the performance of many models!*

# YOUR TURN!

*Using the unnested glance data:*

- 1. Can you find the country models with the highest adjusted  $R^2$ ? What about the lowest?*
- 2. Plot the adjusted  $R^2$  against each continent? What do you find?*
- 3. Filter for adjusted  $R^2 < 0.25$ . What countries do you find? What do you think is driving this bad fit? (Hint: plot the life expectancy over time for these countries)*


# SOLUTION

```
# Can you find the country models with the highest adjusted R2? What about the lowest?
by_country %>%
  mutate(glance = map(model, broom::glance)) %>%
  unnest(glance, .drop = TRUE) %>%
  arrange(adj.r.squared)
# A tibble: 142 × 13
```

	country	continent	r.squared	adj.r.squared	sigma	statistic	p.value
	<fctr>	<fctr>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>
1	Rwanda	Africa	0.01715964	-0.081124401	6.558269	0.1745923	0.68489269
2	Botswana	Africa	0.03402340	-0.062574259	6.112177	0.3522177	0.56604135
3	Zimbabwe	Africa	0.05623196	-0.038144842	7.205431	0.5958240	0.45802898
4	Zambia	Africa	0.05983644	-0.034179918	4.528713	0.6364471	0.44353178
5	Swaziland	Africa	0.06821087	-0.024968046	6.644091	0.7320419	0.41225300
6	Lesotho	Africa	0.08485635	-0.006658011	5.933934	0.9272463	0.35828637
7	Cote d'Ivoire	Africa	0.28337240	0.211709644	3.925590	3.9542491	0.07480350

# SOLUTION

```
# Plot the adjusted R2 against each continent? What do you find?  
by_country %>%  
  mutate(glance = map(model, broom::glance)) %>%  
  unnest(glance, .drop = TRUE) %>%  
  ggplot(aes(continent, adj.r.squared)) +  
  geom_jitter(width = 0.25)
```



# SOLUTION

```
# Filter for adjusted R2 < 0.25.  What countries do you find?
```

```
bad_fit <- by_country %>%  
  mutate(glance = map(model, broom::glance)) %>%  
  unnest(glance, .drop = TRUE) %>%  
  filter(r.squared < 0.25)
```

```
bad_fit
```


```
# A tibble: 6 × 13
```

	country	continent	r.squared	adj.r.squared	sigma	statistic	p.value	df
	<fctr>	<fctr>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<int>
1	Botswana	Africa	0.03402340	-0.062574259	6.112177	0.3522177	0.5660414	2
2	Lesotho	Africa	0.08485635	-0.006658011	5.933934	0.9272463	0.3582864	2
3	Rwanda	Africa	0.01715964	-0.081124401	6.558269	0.1745923	0.6848927	2
4	Swaziland	Africa	0.06821087	-0.024968046	6.644091	0.7320419	0.4122530	2
5	Zambia	Africa	0.05983644	-0.034179918	4.528713	0.6364471	0.4435318	2



# SOLUTION

```
# plot the life expectancy over time for these countries
gapminder %>%
  semi_join(bad_fit, by = "country") %>%
  ggplot(aes(year, lifeExp, colour = country)) +
  geom_line()
```



*likely the tragedies of the  
HIV/AIDS epidemic and  
the Rwandan genocide*






SO LITTLE TIME!





# LEARN MORE



WHAT TO REMEMBER



# FUNCTIONS TO REMEMBER

Operator/Function	Description
<code>nest</code>	Create a nested data frame with list-columns
<code>map2</code>	Similar to the <code>map</code> but will map a specified function over two data inputs
<code>unnest</code>	Unnest our data
<code>broom::glance</code>	Extract model quality metrics into a tidy data structure